

Biconductor Exercises

Gopuraja Dharmalingam

04 February 2015

1. Print few gene names from org.Hs.eg.db

```
library("org.Hs.eg.db")
columns(org.Hs.eg.db)
## [1] "ENTREZID"      "PFAM"          "IPI"           "PROSITE"
## [5] "ACCNUM"        "ALIAS"         "CHR"           "CHRLOC"
## [9] "CHRLOCEND"     "ENZYME"        "MAP"           "PATH"
## [13] "PMID"          "REFSEQ"        "SYMBOL"        "UNIGENE"
## [17] "ENSEMBL"       "ENSEMBLPROT"   "ENSEMBLTRANS"  "GENENAME"
## [21] "UNIPROT"       "GO"            "EVIDENCE"      "ONTOLOGY"
## [25] "GOALL"         "EVIDENCEALL"   "ONTOLOGYALL"   "OMIM"
## [29] "UCSCKG"
head(keys(org.Hs.eg.db, keytype="GENENAME"))
## [1] "alpha-1-B glycoprotein"
## [2] "alpha-2-macroglobulin"
## [3] "alpha-2-macroglobulin pseudogene 1"
## [4] "N-acetyltransferase 1 (arylamine N-acetyltransferase)"
## [5] "N-acetyltransferase 2 (arylamine N-acetyltransferase)"
## [6] "N-acetyltransferase pseudogene"
```

2. Print non-redundant list of chromosomes from org.Mm.eg.db

```
library("org.Mm.eg.db")
unique(keys(org.Mm.eg.db, keytype="CHR"))
## [1] "6" "11" "4" "3" "2" "X" "17" "1" "7" "5" "15" "12" "9" "8"
## [15] "16" "13" "19" "14" "10" "18" "Y" "MT" "Un"
```

3. Retrieve gene name, chromosome and Ensembl gene identifiers for “HEBP2” and “PRND” from org.Hs.eg.db

```
select(org.Hs.eg.db, keys=c("HEBP2", "PRND"), keytype="SYMBOL",
       columns=c("GENENAME", "CHR", "ENSEMBL"))
## SYMBOL GENENAME CHR ENSEMBL
## 1 HEBP2 heme binding protein 2 6 ENSG000000051620
## 2 PRND prion protein 2 (dublet) 20 ENSG00000171864
```

4. Retrieve gene symbol, gene name and gene alias for genes in chromosome 2 for human using org.Hs.eg.db

```
chr2Genes <- select(org.Hs.eg.db,keys="2",keytype="CHR",
                    columns=c("SYMBOL","GENENAME","ALIAS"))
head(chr2Genes)
##   CHR SYMBOL                                GENENAME ALIAS
## 1    2  AAMP  angio-associated, migratory cell protein  AAMP
## 2    2 ACADL  acyl-CoA dehydrogenase, long chain ACAD4
## 3    2 ACADL  acyl-CoA dehydrogenase, long chain LCAD
## 4    2 ACADL  acyl-CoA dehydrogenase, long chain ACADL
## 5    2  ACP1  acid phosphatase 1, soluble HAAP
## 6    2  ACP1  acid phosphatase 1, soluble ACP1
```

5. Retrieve genomic coordinates for human protein coding genes from Ensembl biomart and build GRanges object. Include genes in only main chromosomes (1-22,X,Y).

Tips:

- You can select main chromosomes and “protein coding” genes by using appropriate filter and value.
- Search for “biotype” in available filters using `grep()`
- Run `filterOptions("biotype",selectedmart)` to see the accepted values for “biotype” filter
- When multiple filters specified, “values” argument should be a list of vectors; each vector corresponds to each specified filter.
- Annotation fields to retrieve: “chromosome_name”, “start_position”, “end_position”, “ensembl_gene_id”, “strand”, “external_gene_name”
- Before creating GRanges object, add “chr” prefix to chromosome using `paste` function, Ex: change 1 to chr1 (required for next task)

```
library("biomaRt")
ensembl <- useMart("ensembl") # select ensembl
ens_datasets <- listDatasets(ensembl) # list datasets
ens_human <- useDataset("hsapiens_gene_ensembl",mart=ensembl) # select human dataset
ens_human_attr <- listAttributes(ens_human) # list available annotation
ens_human_filters <- listFilters(ens_human) # list available filters
availFilters <- filterOptions("biotype",ens_human) # Displays accepted values for "biotype"

hg19Gene <- getBM(
  attributes = c("chromosome_name","start_position","end_position",
                "ensembl_gene_id","strand","external_gene_name"),
  filter=c("chromosome_name","biotype"),
  values=list(c(1:22,"X","Y"),"protein_coding"), mart=ens_human)
head(hg19Gene)
##   chromosome_name start_position end_position ensembl_gene_id strand
## 1              20      31465506      31476757 ENSG00000180383     -1
## 2               1       9997206      10016020 ENSG00000162444      1
## 3              X      48186220      48277578 ENSG00000165583     -1
## 4               8      30156297      30183640 ENSG00000104671      1
## 5              20      33407955      33443892 ENSG00000101400     -1
## 6              X      51490011      51496596 ENSG00000196368     -1
##   external_gene_name
## 1          DEFB124
## 2           RBP7
```

```
## 3          SSX5
## 4          DCTN6
## 5          SNTA1
## 6          NUDT11
```

Now create GRanges object using the above data frame.

```
library(GenomicRanges)

# add 'chr' prefix to chromosome name
hg19Gene$chromosome_name <- paste("chr",hg19Gene$chromosome_name,sep="")

hg19Gene.GR <- GRanges(seqnames=hg19Gene$chromosome_name,
                      ranges=IRanges(start=hg19Gene$start_position,end=hg19Gene$end_position),
                      strand=ifelse(hg19Gene$strand==1,"+","-"),
                      EnsemblID=hg19Gene$ensembl_gene_id,
                      Symbol=hg19Gene$external_gene_name)

hg19Gene.GR
## GRanges object with 19850 ranges and 2 metadata columns:
##           seqnames          ranges strand |           EnsemblID
##           <Rle>          <IRanges> <Rle> |           <character>
##      [1]    chr20  [31465506, 31476757]   - | ENSG00000180383
##      [2]    chr1   [ 9997206, 10016020]   + | ENSG00000162444
##      [3]    chrX   [48186220, 48277578]   - | ENSG00000165583
##      [4]    chr8   [30156297, 30183640]   + | ENSG00000104671
##      [5]    chr20  [33407955, 33443892]   - | ENSG00000101400
##      ...      ...      ...      ...      ...
## [19846]    chr1  [171248471, 171285978]   + | ENSG00000010932
## [19847]   chr19 [ 43872363, 43901385]   - | ENSG00000176222
## [19848]    chr1 [ 99708703, 99766631]   - | ENSG00000156869
## [19849]   chr19 [ 57876765, 57916591]   - | ENSG00000269476
## [19850]   chr19 [ 57803841, 57814913]   - | ENSG00000178935
##           Symbol
##           <character>
##      [1]    DEFB124
##      [2]     RBP7
##      [3]     SSX5
##      [4]     DCTN6
##      [5]     SNTA1
##      ...      ...
## [19846]     FMO1
## [19847]    ZNF404
## [19848]    FRRS1
## [19849] CTD-2583A14.9
## [19850]    ZNF552
## -----
## seqinfo: 24 sequences from an unspecified genome; no seqlengths
```

6. Filter the above GRanges object for genes in chr1:1544000-2371000

```

chr1genes <- hg19Gene.GR[seqnames(hg19Gene.GR)=="chr1" &
                        start(hg19Gene.GR) > 1544000 &
                        end(hg19Gene.GR) < 2371000]

head(chr1genes)
## GRanges object with 6 ranges and 2 metadata columns:
##      seqnames      ranges strand |      EnsemblID      Symbol
##      <Rle>        <IRanges> <Rle> |      <character> <character>
## [1]   chr1 [2184461, 2212720]   - | ENSG00000162585   C1orf86
## [2]   chr1 [2019324, 2030751]   + | ENSG00000187730   GABRD
## [3]   chr1 [2050470, 2185395]   + | ENSG00000067606   PRKCZ
## [4]   chr1 [1702730, 1724324]   - | ENSG00000008128   CDK11A
## [5]   chr1 [1615415, 1630610]   + | ENSG00000197530   MIB2
## [6]   chr1 [1598011, 1600096]   - | ENSG00000228594   C1orf233
## -----
##      seqinfo: 24 sequences from an unspecified genome; no seqlengths

# alternate
chr1genes <- subset(hg19Gene.GR,start>1544000 & end<2371000 & seqnames=="chr1")

```

7. Create a GRanges of Transcription start sites (1 bp range) for the GRanges object created in Q5.

Tip:

- How to identify TSS for genes in forward/reverse strand?

```

hg19Gene$TSS <- ifelse(hg19Gene$strand==1,hg19Gene$start_position,hg19Gene$end_position)

hg19TSS <- GRanges(seqnames=hg19Gene$chromosome_name,
                  ranges=IRanges(start=hg19Gene$TSS,end=hg19Gene$TSS),
                  strand=ifelse(hg19Gene$strand==1,"+","-"),
                  EnsemblID=hg19Gene$ensembl_gene_id,
                  Symbol=hg19Gene$external_gene_name)

hg19TSS
## GRanges object with 19850 ranges and 2 metadata columns:
##      seqnames      ranges strand |      EnsemblID
##      <Rle>        <IRanges> <Rle> |      <character>
## [1]   chr20 [31476757, 31476757]   - | ENSG00000180383
## [2]   chr1 [ 9997206,  9997206]   + | ENSG00000162444
## [3]   chrX [48277578, 48277578]   - | ENSG00000165583
## [4]   chr8 [30156297, 30156297]   + | ENSG00000104671
## [5]   chr20 [33443892, 33443892]   - | ENSG00000101400
## ...      ...      ...      ...      ...
## [19846] chr1 [171248471, 171248471]   + | ENSG00000010932
## [19847] chr19 [ 43901385,  43901385]   - | ENSG00000176222
## [19848] chr1 [ 99766631,  99766631]   - | ENSG00000156869
## [19849] chr19 [ 57916591,  57916591]   - | ENSG00000269476
## [19850] chr19 [ 57814913,  57814913]   - | ENSG00000178935
##      Symbol
##      <character>
## [1]      DEFB124

```

```
##      [2]      RBP7
##      [3]      SSX5
##      [4]      DCTN6
##      [5]      SNTA1
##      ...      ...
## [19846]      FM01
## [19847]      ZNF404
## [19848]      FRRS1
## [19849] CTD-2583A14.9
## [19850]      ZNF552
## -----
## seqinfo: 24 sequences from an unspecified genome; no seqlengths
```

8. Create a GRanges object of human promoters with TSS \pm 2000bp (using the GRanges object created in Q5). Tip: Read the documentation for promoters function.

```
hg19Promoters <- promoters(hg19Gene.GR,upstream=2000,downstream=2000)
hg19Promoters
## GRanges object with 19850 ranges and 2 metadata columns:
##      seqnames      ranges strand |      EnsemblID
##      <Rle>      <IRanges> <Rle> |      <character>
##      [1]      chr20 [31474758, 31478757] - | ENSG000000180383
##      [2]      chr1 [ 9995206, 9999205] + | ENSG000000162444
##      [3]      chrX [48275579, 48279578] - | ENSG000000165583
##      [4]      chr8 [30154297, 30158296] + | ENSG000000104671
##      [5]      chr20 [33441893, 33445892] - | ENSG000000101400
##      ...      ...      ...      ...      ...
## [19846]      chr1 [171246471, 171250470] + | ENSG000000010932
## [19847]      chr19 [ 43899386, 43903385] - | ENSG000000176222
## [19848]      chr1 [ 99764632, 99768631] - | ENSG000000156869
## [19849]      chr19 [ 57914592, 57918591] - | ENSG000000269476
## [19850]      chr19 [ 57812914, 57816913] - | ENSG000000178935
##      Symbol
##      <character>
##      [1]      DEFB124
##      [2]      RBP7
##      [3]      SSX5
##      [4]      DCTN6
##      [5]      SNTA1
##      ...      ...
## [19846]      FM01
## [19847]      ZNF404
## [19848]      FRRS1
## [19849] CTD-2583A14.9
## [19850]      ZNF552
## -----
## seqinfo: 24 sequences from an unspecified genome; no seqlengths
```

9. Import ELF1 binding sites in K562 cell from Encode (ELF1_K562.bed) and create GRanges object.

Tips:

- Import the ELF1 binding sites using `import.bed()` function from `rtracklayer` package and compare it with the above GRanges object
- Find ELF1 binding sites overlap with promoters (TSS \pm 1kb) using `findOverlaps` and `subsetByOverlaps`
- Remember BED format uses 0-based coordinates

```
library("rtracklayer")
ELF1 <- read.table("ELF1_K562.bed", sep="\t", header=F)
ELF1GR <- GRanges(seqnames=ELF1$V1, IRanges(start=ELF1$V2+1, end=ELF1$V3))
ELF1GR_A <- import.bed("ELF1_K562.bed")

# ELF1 binding sites overlap with promoters using `findOverlaps`
hg19Promoters <- promoters(hg19Gene.GR, upstream=1000, downstream=1000)
ELF1GR <- reduce(ELF1GR) # merging overlapping peaks
ELF1overlap <- findOverlaps(ELF1GR, hg19Promoters, ignore.strand=T)
ELF1overlap.m <- as.matrix(ELF1overlap)
ELF1_promoters <- ELF1GR[ELF1overlap.m[, "queryHits"],]
ELF1_promoters
## GRanges object with 1480 ranges and 0 metadata columns:
##           seqnames           ranges strand
##           <Rle>             <IRanges> <Rle>
##      [1]      chr1      [ 999543,   999772]      *
##      [2]      chr1     [1115885,  1116069]      *
##      [3]      chr1     [1174741,  1175016]      *
##      [4]      chr1     [1307428,  1307703]      *
##      [5]      chr1     [1310482,  1311094]      *
##      ...      ...      ...      ...
##    [1476]    chrX [130110006, 130110281]      *
##    [1477]    chrX [153763052, 153763327]      *
##    [1478]    chrX [153775333, 153775532]      *
##    [1479]    chrX [153775672, 153775982]      *
##    [1480]    chrX [153777173, 153777448]      *
##    -----
##    seqinfo: 23 sequences from an unspecified genome; no seqlengths

# ELF1 binding sites overlap with promoters using `subsetByOverlaps`
ELF1_promoters1 <- subsetByOverlaps(ELF1GR, hg19Promoters)
ELF1_promoters1
## GRanges object with 1303 ranges and 0 metadata columns:
##           seqnames           ranges strand
##           <Rle>             <IRanges> <Rle>
##      [1]      chr1      [ 999543,   999772]      *
##      [2]      chr1     [1115885,  1116069]      *
##      [3]      chr1     [1174741,  1175016]      *
##      [4]      chr1     [1307428,  1307703]      *
##      [5]      chr1     [1310482,  1311094]      *
##      ...      ...      ...      ...
```

```
## [1299] chrX [130110006, 130110281] *
## [1300] chrX [153763052, 153763327] *
## [1301] chrX [153775333, 153775532] *
## [1302] chrX [153775672, 153775982] *
## [1303] chrX [153777173, 153777448] *
## -----
## seqinfo: 23 sequences from an unspecified genome; no seqlengths
```

Note the differences in the outputs!

10. Retrieve the transcript coordinates for genes as GRangesList from TxDb.Hsapiens.UCSC.hg19.knownGene (install it from Bioconductor if required)

```
# source("http://bioconductor.org/biocLite.R")
# biocLite("TxDb.Hsapiens.UCSC.hg19.knownGene")
library("TxDb.Hsapiens.UCSC.hg19.knownGene")
hg19txdb <- TxDb.Hsapiens.UCSC.hg19.knownGene
TranscriptsByGene <- transcriptsBy(hg19txdb,by="gene") # inspect the output
```

11. Retrieve 200 bp upstream promoter sequences for the given gene symbols AQP1, ASNSP2, KPNA2, FRMD4A, NSUN5, VAC14 from Ensembl human biomart

Tips:

- Read documentation for `getSequence`
- Use `type="hgnc_symbol"` and `seqType="coding_gene_flank"`

```
symbols <- c("AQP1", "ASNSP2", "KPNA2", "FRMD4A", "NSUN5", "VAC14")
seq = getSequence(id=symbols,
                  type="hgnc_symbol",
                  seqType="coding_gene_flank",upstream=200, mart = ens_human)
```

Additional Exercise

- Download the following BAM and index files (*.bai) (ENCODE data - ChIP-Seq of CTCF in Ag04449 human fibroblast cells)
 - <http://hgdownload.cse.ucsc.edu/goldenPath/hg19/encodeDCC/wgEncodeUwTfbs/wgEncodeUwTfbsAg04449CtcfSbam>
 - <http://hgdownload.cse.ucsc.edu/goldenPath/hg19/encodeDCC/wgEncodeUwTfbs/wgEncodeUwTfbsAg04449CtcfSbam.bai>
- Use `readGAlignments` to read the bam. Construct an `ScanBamParam` object that accepts only aligned reads, passing quality control and not duplicates.
- Compute genome wide coverage using `coverage` function
- Compute number of reads overlapping with hg19 promoters (TSS \pm 1kb) and export the results as text file.

```

library("GenomicAlignments")
bamFile <- "wgEncodeUwTfbsAg04449CtcfStdAlnRep1.bam"
flag <- scanBamFlag()
param <- ScanBamParam(
  flag=scanBamFlag(isUnmappedQuery=FALSE, isDuplicate=FALSE, isNotPassingQualityControls=FALSE)
)

# Read the BAM
CTCF <- readGAlignments(bamFile,param=param)

# Generate the genomic coverage and inspect the output
CTFCov <- coverage(CTCF)

# Reads overlapping with promoters
CTCFCounts <- countOverlaps(hg19Promoters,CTCF)

# Add CTCF counts as elementMetadata to hg19Promoters object
mcols(hg19Promoters)$CTCF <- CTCFCounts

# Export the results as text file
hg19Promoters.df <- as.data.frame(hg19Promoters)
write.table(hg19Promoters.df,"hg19Promoters_CTCF.txt",sep="\t",row.names=F)

```